

Memory Management



Binding of Instructions and Data to Memory

Address binding of instructions and data to memory addresses can happen at three different stages

- **Compile time:** If memory location known a priori, *absolute code* can be generated; must recompile code if starting location changes
- **Load time:** Must generate *relocatable code* if memory location is not known at compile time
- **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another. Need hardware support for address maps (e.g., *base* and *limit registers*).

8.4

Memory Management

- Background
- Swapping
- Contiguous Allocation
- Paging
- Segmentation
- Segmentation with Paging

8.2

Logical vs. Physical Address Space

- The concept of a logical *address space* that is bound to a separate *physical address space* is central to proper memory management
 - Logical address** – generated by the CPU; also referred to as *virtual address*
 - Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme

8.5

Background

- Program must be brought into memory and placed within a process for it to be run
- **Input queue** – collection of processes on the disk that are waiting to be brought into memory to run the program
- User programs go through several steps before being run


8.3

Memory-Management Unit (MMU)

- Hardware device that maps virtual to physical address
- In MMU scheme, the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- The user program deals with *logical* addresses; it never sees the *real* physical addresses

8.6

Dynamic relocation using a relocation register



8.7

7

Swapping

- A process can be swapped temporarily out of memory to a backing store, and then brought back into memory for continued execution
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)

8.10

10


Dynamic Loading

- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required implemented through program design

8.8

8

Schematic View of Swapping



8.11

11

Dynamic Linking

- Linking postponed until execution time
- Small piece of code, *stub*, used to locate the appropriate memory-resident library routine
- Stub replaces itself with the address of the routine, and executes the routine
- Operating system needed to check if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries

8.9

9


Contiguous Allocation

- Main memory usually into two partitions:
Resident operating system, usually held in low memory with interrupt vector
User processes then held in high memory
- Single-partition allocation
Relocation-register scheme used to protect user processes from each other, and from changing operating-system code and data
Relocation register contains value of smallest physical address; limit register contains range of logical addresses – each logical address must be less than the limit register

8.12

12

A base and a limit register define a logical address space



8.13


13

Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes

- First-fit: Allocate the *first* hole that is big enough
- Best-fit: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size. Produces the smallest leftover hole.
- Worst-fit: Allocate the *largest* hole; must also search entire list. Produces the largest leftover hole.

First-fit and best-fit better than worst-fit in terms of speed and storage utilization




Example Memory Configuration Before and After Allocation of 16 Kbyte Block

8.16

16

HW address protection with base and limit registers




8.14

14

Fragmentation

- External Fragmentation – total memory space exists to satisfy a request, but it is not contiguous
- Internal Fragmentation – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- Reduce external fragmentation by compaction




8.17

17

Contiguous Allocation (Cont.)

- Multiple-partition allocation

Hole – block of available memory; holes of various size are scattered throughout memory
When a process arrives, it is allocated memory from a hole large enough to accommodate it
Operating system maintains information about:
a) allocated partitions b) free partitions (hole)



8.18

15

Paging

- Logical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
- Divide physical memory into fixed-sized blocks called **frames** (size is power of 2, between 512 bytes and 8192 bytes)
- Divide logical memory into blocks of same size called **pages**.
- Keep track of all free frames
- To run a program of size n pages, need to find n free frames and load program
- Set up a page table to translate logical to physical addresses
- Internal fragmentation

8.19

18

Address Translation Scheme


- Address generated by CPU is divided into:

Page number (p) – used as an index into a **page table** which contains base address of each page in physical memory

Page offset (d) – combined with base address to define the physical memory address that is sent to the memory unit


8.19

Paging Example




8.22

Address Translation Architecture




8.20

Free Frames



8.23

Paging Example



8.21

Implementation of Page Table

- Page table is kept in main memory
- Page-table base register (PTBR)** points to the page table
- Page-table length register (PRLR)** indicates size of the page table
- In this scheme every data/instruction access requires two memory accesses. One for the page table and one for the data/instruction.
- The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**

8.24

Associative Memory

- Associative memory – parallel search

Page #	Frame #

Address translation (A' , A'')

If A' is in associative register, get frame # out

Otherwise get frame # from page table in memory


8.25

Memory Protection

- Memory protection implemented by associating protection bit with each frame
- Valid-invalid** bit attached to each entry in the page table:
 - "valid" indicates that the associated page is in the process' logical address space, and is thus a legal page
 - "invalid" indicates that the page is not in the process' logical address space

8.28

Paging Hardware With TLB



8.26

Valid (v) or Invalid (i) Bit In A Page Table

00000	page 0	0	2	v	0
	page 1	1	3	v	1
	page 2	2	4	v	2
	page 3	3	7	v	3
	page 4	4	8	v	4
	page 5	5	9	v	5
10,468		6	0	i	
12,287		7	0	i	

8.29

Effective Access Time

- Associative Lookup = c time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio – percentage of times that a page number is found in the associative registers; ration related to number of associative registers
- Hit ratio = α
- Effective Access Time (EAT)**

$$EAT = (1 + \alpha) \alpha + (2 + \alpha)(1 - \alpha)$$

$$= 2 + \alpha - \alpha$$

8.27

Page Table Structure

- Hierarchical Paging
- Hashed Page Tables
- Inverted Page Tables

8.30

Hierarchical Page Tables


- Break up the logical address space into multiple page tables
- A simple technique is a two-level page table

8.31

31

Address-Translation Scheme

- Address-translation scheme for a two-level 32-bit paging architecture



8.34

34

Two-Level Paging Example

- A logical address (on 32-bit machine with 4K page size) is divided into:
 - a page number consisting of 20 bits
 - a page offset consisting of 12 bits
- Since the page table is paged, the page number is further divided into:
 - a 10-bit page number
 - a 10-bit page offset
- Thus, a logical address is as follows:

page number		page offset
p_1	p_2	d
10	10	12

where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the outer page table

8.32

32


Inverted Page Table

- One entry for each real page of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- Use hash table to limit the search to one — or at most a few — page-table entries

8.35

35


Two-Level Page-Table Scheme



8.33

33

Inverted Page Table Architecture



8.36

36

Shared Pages

Shared code

One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

Shared code must appear in same location in the logical address space of all processes


Private code and data

Each process keeps a separate copy of the code and data

The pages for the private code and data can appear anywhere in the logical address space

8.37

User's View of a Program




8.40

37


40

Shared Pages Example



8.38

Logical View of Segmentation



8.41

38

41

Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments. A segment is a logical unit such as:
 - main program,
 - procedure,
 - function,
 - method,
 - object,
 - local variables, global variables,
 - common block,
 - stack,
 - symbol table, arrays

8.39

Segmentation Architecture

- Logical address consists of a two tuple:
 - <segment-number, offset>.
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - base* – contains the starting physical address where the segments reside in memory
 - limit* – specifies the length of the segment
- **Segment-table base register (STBR)** points to the segment table's location in memory
- **Segment-table length register (STLR)** indicates number of segments used by a program;
 - segment number s is legal if $s < STLR$

8.42

39


42

Segmentation Architecture (Cont.)

- **Relocation.**
dynamic
by segment table
- **Sharing.**
shared segments
same segment number
- **Allocation.**
first fit/best fit
external fragmentation

8.43

Example of Segmentation




8.46

Segmentation Architecture (Cont.)

- **Protection.** With each entry in segment table associate:
validation bit = 0 \Rightarrow illegal segment
read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram


8.44

Sharing of Segments



8.47

Address Translation Architecture



8.48

Segmentation with Paging – MULTICS

- The MULTICS system solved problems of external fragmentation and lengthy search times by paging the segments
- Solution differs from pure segmentation in that the segment-table entry contains not the base address of the segment, but rather the base address of a *page table* for this segment

8.49

